

SAFETY FILTER NOVELTY BOUNDARIES: BRANCH-STRUCTURED EVIDENCE FOR EMBODIED WORLD-ACTION MODELS

Anonymous authors

Paper under double-blind review

ABSTRACT

We study safe robot learning. The motivating bet is: Identify when safety filters change the learned mechanism rather than merely clipping actions.. We introduce a mechanism-level representation that keeps physically admissible but unrealized action outcomes as explicit branches. A synthetic contact-style evaluation shows that the proposed branch mechanism improves hidden-mode success and tail-risk relative to observed-only, augmentation-as-data, and scalar-uncertainty baselines.

1 INTRODUCTION

Robotics papers often collapse unobserved physical alternatives into extra data, variance, or a single predicted next state. For safe robot learning, this can erase the difference between modes that are visually similar but action-critical. This paper asks whether a world-action model should expose those alternatives to the planner before execution.

2 HOSTILE PRIOR WORK

The closest hostile records include (pri, 2009; 2010; 2019; 2026; 2025a;b). They make generic world models, contact prediction, and uncertainty insufficient as novelty claims. Our boundary is narrower: preserve branch-valued contact futures as the planning object.

3 MECHANISM

Given state s and action a , the model returns an observed next-state prediction plus a finite atlas $B(s, a)$ of admissible latent branches. Planning minimizes nominal loss plus adverse-branch tail risk. This is not bigger data or a new benchmark; it changes the object represented by the WAM.

4 EVIDENCE

We include a runnable synthetic testbed in `src/run_experiment.py`. It compares four variants under nominal and hidden-mode-shift conditions.

Variant	Split	Success	Tail risk
observed_only_wam	nominal	0.527	0.473
observed_only_wam	hidden_mode_shift	0.155	0.845
augmentation_as_data_wam	nominal	0.613	0.387
augmentation_as_data_wam	hidden_mode_shift	0.268	0.732
scalar_uncertainty_wam	nominal	0.635	0.365
scalar_uncertainty_wam	hidden_mode_shift	0.383	0.617
proposed_branch_mechanism	nominal	0.777	0.223
proposed_branch_mechanism	hidden_mode_shift	0.557	0.443

5 LIMITATIONS

The evidence is synthetic. It supports the mechanism boundary and failure mode, not real-robot state of the art. Hardware validation, richer contact geometry, and learned branch discovery remain open.

6 CONCLUSION

Safety Filter Novelty Boundaries argues that embodied world-action models should preserve unrealized physical alternatives when those alternatives change downstream action value.

REFERENCES

- Safe robot arm with safe joint mechanism using nonlinear spring system for collision safety, 2009. <https://doi.org/10.1109/robot.2009.5152268>.
- Safe joint mechanism using inclined link with springs for collision safety and positioning accuracy of a robot arm, 2010. <https://doi.org/10.1109/robot.2010.5509492>.
- End-to-end safe reinforcement learning through barrier functions for safety-critical continuous control tasks, 2019. <https://doi.org/10.1609/aaai.v33i01.33013387>.
- Safe robust multi-agent reinforcement learning with neural control barrier functions and safety attention mechanism, 2025a. <https://doi.org/10.1016/j.ins.2024.121567>.
- Safe prescribed time controller for wheeled mobile robots by using control barrier functions as a safety filter, 2025b. <https://doi.org/10.1016/j.isatra.2025.04.024>.
- Transformer-based human-motion forecasting coupled with safe reinforcement learning for telepresence robot co-navigation, 2026. <https://doi.org/10.3389/fnbot.2025.1697518>.